

431 Project Instructions

Thomas E. Love

Version: 2018-09-16 21:16:19

Contents

Overview	5
Your Project includes Two Studies	5
You have Nine Tasks to Complete this Semester	5
Working with This Document	5
Need Help?	6
1 Project Objectives	7
1.1 Study 1 is about making comparisons and visualizing groups of data.	7
1.2 Study 2 is about building a model, and making predictions.	7
1.3 Why Two Studies?	8
1.4 Educational Objectives	8
2 Task A (The Proposal) Instructions	11
2.1 Deadline and Submission information	11
2.2 Study 1 work for Task A	11
2.3 Study 2 Work for Task A	15
2.4 Evaluating Task A	18
3 Task B (Presentation Sign-Up) Instructions	19
3.1 Deadline and Submission information	19
4 Task C (Survey Editing) Instructions	21
4.1 Deadline and Submission information	21
5 Task D (Survey Comparison Plan) Instructions	23
5.1 Deadline and Submission information	23
6 Task E (Taking the Survey) Instructions	25
6.1 Deadline and Submission information	25
7 Task F (Sharing Study 2 Data) Instructions	27
7.1 Deadline and Submission information	27
7.2 Sharing Your Data Appropriately	27
8 Task G (The Project Update) Instructions	29
8.1 Deadline and Submission information	29
9 Task H (The Portfolio) Instructions	31
10 Task I (Your Presentation) Instructions	33

Overview

This website contains the Fall 2018 project information for PQHS / CRSP / MPHP 431: Statistical Methods in Biological & Medical Sciences, Section 1.

- All materials related to the project (including these instructions) are maintained and linked at <https://github.com/THOMASELOVE/431-2018-project>.
- The direct link to this document is <https://thomaselove.github.io/431-2018-project>.

Your Project includes Two Studies

Your final project for this course will result in a portfolio of work related to two studies.

Study 1 - Class Survey. In the first study, you (sometimes working individually, sometimes in a group) will design, administer, analyze and present the results of a survey designed to compare two or three groups of subjects on some *categorical* and *quantitative* outcomes we will develop from your initial ideas.

Study 2 - Your Data. In the second study, you (working individually) will propose a research question and relevant data of interest to you, and then complete all elements of a data science project designed to create a statistical model for a *quantitative* outcome, then use it for prediction and assess the quality of those predictions.

You have Nine Tasks to Complete this Semester

The project involves two analyses (one for the class survey and one for your personal study), and a total of 9 tasks (deliverables.) Each task is to be completed by **12 NOON** on the specified date.

- Task A (The Proposal) is due at noon on 2018-10-12
- Task B (Presentation Sign-Up) is **also** due at noon on 2018-10-12
- Task C (Survey Editing) involves group work and is due at noon on 2018-10-23
- Task D (Survey Comparison Plan) is **also** due at noon on 2018-10-23
- Task E (Taking the Survey) is due at noon on 2018-10-31
- Task F (Sharing Study 2 Data) is due at noon on 2018-11-14
- Task G (The Update) is due at noon on 2018-11-28
- Task H (The Portfolio) is due at noon on 2018-12-13
- Task I (Your Presentation) will be held on 2018-12-10, 2018-12-11 or 2018-12-13

The bulk of this document contains specific instructions for each of these tasks.

Working with This Document

1. This document is broken down into multiple sections. Use the table of contents at left to navigate.
2. At the top of the document, you'll see icons which you can click to

- search the document,
 - change the size, font or color scheme of the page, and
 - download a PDF or EPUB (Kindle-readable) version of the entire document.
3. The document is a work in progress, and will be updated occasionally through the semester. Check the Version information above to verify the last update time¹.

Need Help?

Questions about the project or the course can be directed to **431-help at case dot edu** or to Dr. Love directly at **thomas dot love at case dot edu**.

- The course home page is at <https://github.com/THOMASELOVE/431-2018>

¹Note that the ePub and PDF versions will show slightly different times (but on the same day) as the HTML version.

Chapter 1

Project Objectives

It is hard to learn statistics (or anything else) passively; concurrent theory and application are essential¹.

1.1 Study 1 is about making comparisons and visualizing groups of data.

Study 1 involves data from a **class survey**, to be conducted in October. We will design, administer, analyze and present survey results designed to compare two or three groups of subjects from the class on some *categorical* and *quantitative* outcomes. In the analysis stage, everyone will be working with different parts of the same data set.

Think of a graph as a comparison. All graphs are comparisons (indeed, all statistical analyses are comparisons). If you already have the graph in mind, think of what comparisons it's enabling. Or if you haven't settled on the graph yet, think of what comparisons you'd like to make. Andrew Gelman

In your eventual analysis of Study 1, you will be comparing both quantitative and categorical outcomes across 2-3 groups. All tools necessary for Study 1 are in Parts A and B of the course, and include the following...

- Descriptive and exploratory summaries of the data across the groups for each of your chosen outcomes, including, of course, attractive and well-constructed visualizations, graphs and tables.
- Comparisons of the population mean difference for at least one quantitative outcome across a set of two (or three) groups, including appropriate demonstrations of the reasons behind the choices you made between parametric, non-parametric and bootstrap procedures.
- Comparisons of the population proportions for at least one categorical outcome across your set of two (or three) groups, including appropriately interpreted point estimates and confidence intervals.

Note well that Study 1 is **not** about building sophisticated statistical models, and using them to make predictions. That's Study 2.

1.2 Study 2 is about building a model, and making predictions.

Study 2 involves data about a **research question that you will propose**, involving data of interest to you. Thus, everyone will be working with a different data set. You will complete all elements of a data

¹Though by no means an original idea, this particular phrasing is stolen from Harry Roberts.

science project designed to create a statistical model for a *quantitative* outcome, then use it for prediction, and assess the quality of those predictions.

All models are wrong but some are useful. George E. P. Box

In Study 2, you will be building a multiple linear regression model, and using it to predict a quantitative outcome of interest. The tools necessary for Study 2 appear in each Part of the course, especially Part C, and include the following...

- Describing the experimental or observational study design used to gather the data, as well as the complete data collection process.
- Sharing the complete raw data in an appropriate way with a statistician (Dr. Love). This means that, in general, data including protected health information are *not* appropriate for this project.
- Developing appropriate research questions that lead to the identification of smart measures for predictors and outcomes, and then the development of a prediction model using multiple linear regression.
- Using a training sample to develop a model, and present the process that leads to a final set of 2-3 candidate models in the training sample.
- Using a test sample to evaluate the quality of predictions from each of the candidate models, and making a final selection.
- Evaluating the adherence of the data you’ve collected to the assumptions of multiple linear regression, and iterating through the model-building process as necessary until the final model shows no strong violations of those assumptions.

1.3 Why Two Studies?

The main reason is that I can’t figure out a way to get you to think about all of the things I hope you’ll learn from this project in a single study.

1. I set different tasks for Study 1 and for Study 2, allowing us to touch on a wider fraction of the things I hope you’ll learn in 431.
2. I want some of the work to be done as a class, some in groups, some as individuals.
3. Some of you have easy access to great data you want to study in this class, and in fact, that’s a primary motivation for taking the class. But not all of you.
4. I have to evaluate each of your projects, and there are many students in the class. Knowing at least one of the data sets you’ll be working with helps me manage this.
5. Having a broad range of activities to evaluate helps reduce the cost of a mistake on any one of them, so that we can build on what you do well.
6. All of Study 1 can be done by the middle of November, leaving the last few weeks of the semester for you to focus on Study 2.

1.4 Educational Objectives

“Statistics has no reason for existence except as the catalyst for investigation and discovery.”
George E. P. Box

I am primarily interested in your learning something interesting, useful and even valuable from your project. An effective project will demonstrate:

1. The ability to create and formulate research questions that are statistically and scientifically appropriate.
2. The ability to turn research questions into measures of interest.
3. The ability to pull and merge and clean and tidy data, then present the data set following Jeff Leek’s guide to sharing data with a statistician.
4. The ability to identify appropriate estimation / testing procedures for the class survey using both continuous and categorical outcomes.

5. The ability to build a reasonable model, including interactions and transformations to deal with non-linearity, assess the quality of the model and residual plots, then use the model to make predictions.
6. The ability to build a Table 1 to showcase potential differences between variables.
7. The ability to identify and (with help) solve problems that crop up
8. The ability to comment on your work within code, and in written and oral presentation.
9. The ability to build a Markdown-based report and a Markdown-based set of slides for presentation.

Chapter 2

Task A (The Proposal) Instructions

Task A requires you to accomplish the following:

On 2018-09-25, you will become part of a **group** of about five people, and your group will:

1. develop and propose 2-3 “research questions” for Study 1 (The Class Survey). Here, your research questions must clearly identify meaningful statistical comparisons.
2. propose 6-10 “homemade” survey questions for Study 1 that relate to your research questions, and
3. propose a “scale” for Study 1 that also relates to your research questions.

As an **individual**, you will also:

4. develop and propose a meaningful summary of your ideas and research question for Study 2 (Your Data). Your research question needs to clearly relate to modeling and prediction of a quantitative outcome on the basis of a set of predictor variables.
5. identify and present a detailed description of a data set that is likely to lead to an answer to the research question proposed for Study 2, and that is appropriate for use in this project.

The rest of this section contains guidance as to what sort of questions your group will need to propose for the class survey in Study 1, and as to what sorts of data sets and research questions are appropriate for your Study 2 proposal.

2.1 Deadline and Submission information

The project groups will only apply to Project Tasks A, C and D. You will become part of a project group on 2018-09-25 and the groups will disband after the survey is finalized in late October.

Task A is due at noon on 2018-10-12. You will submit your Task A work through Canvas.

- The group work (Parts 1-3) need to be submitted by one of the members of your group.
 - If you are not the person submitting this information for your group, then your Task A submission should begin with the statement: “Parts 1-3 of Task A were submitted for my group by PERSON’S NAME.”
- All students need to submit Parts 4 and 5 of Task A, individually.
- Please note that **Task A** and **Task B** are due at the same time.

2.2 Study 1 work for Task A

In Study 1, you will survey your fellow students through an online instrument (containing somewhere around 100 items) that we will develop in Tasks A, C and D in September and October and then administer in the

final week of October (Task E).

Students in the class will develop the items for this instrument in 10 groups of about 5 people per group. The final survey will include questions generated by each of the 10 groups in the class.

The course survey will be done online, and the respondents (de-identified, of course) will include all students in the current 431 class, plus the teaching assistants, and perhaps some volunteers from previous iterations of the course, in an effort to get a reasonable (but by no means random or representative) sample of graduate students at CWRU.

Remember that Study 1 is about making comparisons between groups.

2.2.1 Research Questions

The first part of Task A requires your group to develop and propose 2-3 “research questions” for Study 1 (The Class Survey).

A research question is the fundamental core of a research project, study, or review of literature. It focuses the study, determines the methodology, and guides all stages of inquiry, analysis, and reporting. Source

- The research questions your project group will prepare for Study 1 should state the study objective in terms that allow us to apply statistical methods to test data to obtain an answer.
- Each research question should be written in the form of a comparison of 2-3 exposures or groups in terms of an outcome.
- At least one of your research questions needs to compare groups on a quantitative outcome, and at least one needs to compare groups on a categorical outcome (containing no more than 5 categories.)

Quoting Roger Peng, from Exploratory Data Analysis with R:

Formulating a question can be a useful way to guide the exploratory data analysis process and to limit the exponential number of paths that can be taken with any sizeable dataset. In particular, a sharp question or hypothesis can serve as a dimension reduction tool that can eliminate variables that are not immediately relevant to the question. For example, (suppose that we are interested in) looking at an air pollution dataset from the U.S. Environmental Protection Agency (EPA).

A general question one could ask is “Are air pollution levels higher on the east coast than on the west coast?” But a more specific question might be “Are hourly ozone levels on average higher in New York City than they are in Los Angeles?”

Note that both questions may be of interest, and neither is right or wrong. But the first question requires looking at all pollutants across the entire east and west coasts, while the second question only requires looking at single pollutant in two cities. It’s usually a good idea to spend a few minutes to figure out what is the question you’re really interested in, and narrow it down to be as specific as possible (without becoming uninteresting).

2.2.1.1 Checklist for Research Questions

1. Is our research question (RQ) something that we are curious about and that others might care about?
 2. Does our RQ present an issue on which we can justify a stand prior to data collection about what we think will happen?
 3. Is our RQ too broad, too narrow, or OK, given the time frame and restrictions of this survey?
 4. Is our RQ measurable? What type of information do we need? Can I find a way to ask a limited number of survey questions in such a way to allow me to (after the data are collected) either support or contradict a position on my RQ?
- Adapted from this online tutorial from Empire State College

2.2.1.2 Tips on Writing Good Research Questions

- Duke has a nice overview online of key issues.
- Vanderbilt has some nice materials, built from the tutorial at Empire State College quoted earlier
- Jeff Leek provides several relevant tips in *The Elements of Data Analytic Style*
- <https://researchrundowns.com/intro/writing-research-questions/> has some excellent tips on wording

2.2.2 Specifying Survey Questions

The second part of Task A requires your group to develop and propose 6-10 “homemade” survey questions for Study 1 that relate to your research questions.

Your group will need to specify the exact wording for your survey questions (and potential answers for any categorical responses.) This will likely require some editing and rework, once we have the complete set of proposed questions from all students. Be prepared to revise and resubmit, quickly, so that all items can be resolved in time for publication of the draft survey.

Of your 6-10 survey questions ...

- at least two should ask the respondents to provide you with a **number** that expresses a quantitative outcome of interest to you, and these outcomes should relate closely to at least one of your research questions.
 - If you are asking people to respond to a prompt using a rating, that rating should be expressed on a wide scale. Our preference is a 0-100 scale for quantitative items, where 0 represents the most negative reaction and 100 the most positive reaction to the item.
 - One common choice is to make a statement and ask for agreement with that statement on a scale from 0 = Strongly Disagree to 100 = Strongly Agree.
 - The reason we prefer a 0-100 scale is to increase variation in our responses, as compared to, say, a 1-10 scale.
 - When responding to items using a scale like this on the actual survey, please use the whole scale.
- **at least two** should ask the respondents to provide you with a response that expresses a categorical outcome of interest to you, and these, too, should relate to at least one of your research questions.
 - You will need to specify each of the available responses that you wish to use in the survey.
 - No more than five options for your categorical outcome, please.
 - The response options you specify should be mutually exclusive and collectively exhaustive.
- **at least two** should ask the respondents to categorize themselves into one of two (or three) groups.
 - Be aware that you will need to have at least 10 students in each group in order to build a semi-reasonable analysis.
 - You should expect that 50-55 people will actually respond to the survey, in total.
 - Again, these groupings should be linked to your research questions.

You are welcome to submit exactly 6, or as many as 10 total survey questions in this part of the Task. It is likely that some of your survey questions will not correspond to some of your research questions, and that’s OK, but each survey question should be linked to at least one of your research questions.

2.2.2.1 Dr. Love will include 15 Survey Questions Automatically

The following items will be included in the survey to identify “groups” of students in a reasonable way. As a result, you should not ask these questions in your proposed list, although you can and should consider whether these groupings would be good candidates for application to your research questions.

The following 7 items will have yes/no responses, and thus produce binary groups for analysis.

1. Were you born in the United States?
2. Is English the language you speak better than any other?
3. Do you identify as female?

4. Do you wear prescription glasses or contact lenses?
5. Before taking 431, had you ever used R before?
6. Are you currently married or in a stable domestic relationship?
7. Have you smoked 100 cigarettes or more in your entire life?

The next eight items generate non-binary responses. Together, after the survey is complete, we will identify “cutpoints” for these eight items to identify groups of meaningful size.

8. In what year were you born?
9. How would you rate your current health overall (Excellent, Very Good, Good, Fair, Poor)
10. For how long, in months, have you lived in Northeast Ohio?
11. What is your height in inches? (If you are five feet, eight inches tall, please write 68 inches. To convert from centimeters to inches, multiply your height in centimeters by 0.3937, and then round the result to the nearest inch.)
12. What is your weight in pounds? (To convert from kilograms to pounds, multiply your weight in kilograms by 2.2046, and then round the result to the nearest pound.)
13. What is your pulse rate, in beats per minute? (Please either use a tracking device, or count your pulse for 15 seconds then multiply by 4)
14. Last week, on how many days did you exercise? (0 - 7)
15. Last night, how many hours of sleep did you get?

2.2.2.2 Permitted Types of Items

The survey will be conducted using a Google Form, rather than Survey Monkey or some other tool. Thus, we have a somewhat restricted set of item types.

For **quantitative measures**, Google Forms permit the use of

1. a *short answer* item without any restrictions on the response (except a character limit)
2. a *short answer* item where respondents are forced to insert a number within a given range through a validation process that only accepts the response if it falls within the specified limits.
3. *linear scale* items for ordered categorical ratings (but only on a scale of up to ten points - i.e. 1 to 10)

For **categorical measures**, Google Forms permit the use of

1. *multiple choice* items for endorsing a single choice from a group of 2-10 alternatives.
2. *checkbox* items for the endorsement of one or more choices from a group of 2-10 alternatives.
3. *linear scale* items for ordered categorical ratings (often on a 1-X scale, where X is between 2 and 10)
4. *dropdown* items for selections of one option from a group of 2-3 choices.

2.2.2.3 Old Class Surveys

The surveys (each containing at least 100 items) from 2014, 2015, 2016 and 2017 are available as PDF documents on our web site.

Remember that the rules used this year have been modified from what has been used for the project previously.

2.2.3 Specifying a “Scale”

The third part of Task A requires your group to identify and propose a “scale” of items for Study 1 (The Class Survey).

Your group needs to specify a published scale (available for free public use) to generate an outcome or grouping(s) of interest from those completing the survey. You will have to provide a complete reference to the scale (online, ideally) and specify each of the items in the scale, and how the scale is then evaluated, in all necessary detail to allow us to review and replicate the scale in practice.

The word “scale” is used in many different ways. In this case, by a scale I mean a published list of items, usually accompanied by a scoring rubric that provides some sort of composite score or scores. Examples of scales we have used in the past include:

- Two Health Consciousness Scales, one due to Gould another to Dutta-Bregman Gould Health Consciousness Scale
- The Ten-Item Personality Inventory
- The Perceived Stress Scale
- The Epworth Sleepiness Scale

Your group will need to verify explicitly in your Task A materials that the scale your group proposes is freely available for use by anyone, without any fees or registration requirements.

2.3 Study 2 Work for Task A

You, individually, will present a proposal **summary** (< 300 words) and a brief **data description** for Study 2 in Task A.

You will be building a multiple linear regression model, and using it to predict your outcome of interest.

2.3.1 The Proposal Summary

The fourth part of Task A is to develop and propose a meaningful summary of your ideas and research question for Study 2 (Your Data).

Your summary should begin with a title for your Study 2. Take the time to come up with a good, interesting title. You are going to work hard on this thing; please resist the temptation to kill my interest at the start by calling it “431 Statistics Project” or anything else that shows a similar lack of effort.

Provide me a very brief summary of what you’re trying to accomplish - specifically, what your research question is, and what you hypothesize will happen.

The five most important things to do in the summary are:

1. Write clearly. My best advice is to finish the summary as soon as you can, and then give it to someone else to read, who can criticize it for lack of clarity in the writing.
2. Specify the topic of interest, and motivate your study of it.
3. Explicitly specify your key research question, which should be stated as a question, and which should clearly and naturally lead to a prediction model for a quantitative outcome.
4. Explain what you hypothesize will happen, and
5. Explicitly link your key research question to the data set you describe in your data description.

The summary is the heart of the proposal, and requires some care. You will need to convince me that your topic is interesting, your data are relevant, and building a model and making predictions of a quantitative outcome using the predictors available to you will be worthwhile.

- The summary ends with a statement of the research question or questions (you may have one, or possibly two.) An excellent question conveys the main objective of the study in terms that allow us to apply statistical models to describe an association between one or more predictors and a quantitative outcome. Some advice on writing a good research question is provided below, after the data set description information.
- It should be possible for me to explain your study accurately just by reading this summary. If it’s not possible, it will come back to you for a REDO. Statistics is a details business. Get the details right.
- This summary should be less than 300 words.
- Use complete English sentences. Write in plain language. Use words we all know. Avoid jargon. And look at the general suggestions about writing in the Course Syllabus.

2.3.2 The Data Description

Your data description can be as long as it needs to be, although two pages is usually more than enough. It should include:

1. Your data source, which can be an online source (in which case include a working link), a published paper or journal article (in which case I need a link and a PDF copy of the paper), or unpublished data (in which case I need the details of how the data were gathered).
2. A thorough description of the data collection process, with complete details as to the nature of the variables, the setting for data collection, and complete details of any apparatus you used which may affect results.
3. Specification of the people and methods involved.
 - Who are the subjects under study?
 - When were the data gathered? By whom?
 - How many subjects are included?
 - What caused subjects to be included or excluded from the study?
4. Your planned **quantitative** outcome, which must relate directly to the research question you specified above. Provide a complete definition, including specifying the exact wording of the question or details of the measurement procedure used to obtain the outcome. If available, you can also include descriptions of secondary **quantitative** outcomes. Your outcomes must be quantitative in Study 2.
5. Your predictors of interest, which should also relate to the research question in an obvious way. Again, define the variables carefully, as you did with the outcome.
6. If you already have the data, tell me that. If you don't, specify any steps you must still take in order to get the data, and specify the date by which you will have your data (must be no later than November 1.)

2.3.3 Data Restrictions

Study 2 data sets MUST

- contain between 250 and 250,000 distinct observations,
- contain at least one quantitative outcome variable,
- contain at least four predictor variables, one of which may be identified as the “key” predictor of interest,
- include at least one quantitative predictor variable, and at least one categorical predictor variable,
- include a complete description of how the data were gathered, so that information must be publicly available,
- be in your hands no later than November 1,
- be shared data with a statistician (Dr. Love) following Jeff Leek's guide to sharing data with a statistician as part of Task F. This means you need to have access to the data in the raw, and it means that I have to be able to have access to it in the raw, as well. - be capable of being fully cited for any and all data elements, including a complete codebook, as this must be provided as part of your proposal.

While there are some great resources available to some people in this class by virtue of their affiliation with one of the health systems in town, I can do nothing to get you access to health system specific data as part of your project for this class or for 432, and in general, data from those sources are not especially appropriate because of issues with protected health information.

No more than two students in the class can work on the same data. If two of you have data you would each like to work on, that may be OK, but you'll need to generate separate research questions and perform your analyses and the Project Tasks separately.

I am not interested in you using pre-cleaned data from an educational repository, such as:

- this one at the Cleveland Clinic, or this one at Vanderbilt University, or this one at UCLA, or this one at the University of Florida, or this one at Florida State University, or
- StatLib at Carnegie-Mellon University, or the Journal of Statistics Education Data Archive, or
- the data sets gathered in the fivethirtyeight package, the mosaic package, the cars package, the datasets package, or any other R package designed primary for teaching, or
- StatSci.org's repository of textbook examples and ready for teaching data, or
- any of the many textbook-linked repositories of data sets, like this one for Statistics: Unlocking the Power of Data, or
- any similar repository Professor Love deems to be inappropriate

2.3.4 Some Potentially Useful Data Sources

The ideal choice of data source for this project is a public-use version of a meaningful data set without access restrictions. With so many students in the class, I cannot be responsible for supervising your work with restricted data personally. Some appealing sources to explore include:

- the new Google Datasets Search
- <https://www.data.gov/> The home of the U.S. Government's open data
- <http://www.census.gov/data.html> The U.S. Census Bureau has many interesting data sets, including the Current Population Survey
- <http://www.healthdata.gov/> 125 years of U.S. Health Care Data
- <http://www.cdc.gov/nchs/nhanes/index.htm> National Health and Nutrition Examination Survey.
 - You may want to look at the nhanesA package in R
- <http://dashboard.healthit.gov/datadashboard/data.php> Office of the National Coordinator for Health IT's dashboard
- <http://www.icpsr.umich.edu/icpsrweb/> ICSPR (Inter-university Consortium for Political and Social Research) is a source for many public-use data sets
 - This includes the Health and Medical Care data archive of the Robert Wood Johnson Foundation
- <http://gss.norc.org/> The General Social Survey
- <http://www.bls.gov/data/> Bureau of Labor Statistics
- <http://nces.ed.gov/surveys/> National Center for Education Statistics
- <http://www.odh.ohio.gov/healthstats/dataandstats.aspx> Ohio Department of Health
- <http://open.canada.ca/en> Canada Open Data
- <http://digital.nhs.uk/home> Health data sets from the UK National Health Service.
- <http://www.who.int/en/> World Health Organization
- <http://www.unicef.org/statistics/> UNICEF has some available data on women and children
- <http://www.pewinternet.org/datasets/> Pew Research Center's Internet Project
- <http://portals.broadinstitute.org/cgi-bin/cancer/datasets.cgi> Broad Institute's Cancer Program
- <http://www.kdnuggets.com/datasets/index.html> is a big index of lots of available data repositories
- <https://www.kaggle.com/> Kaggle competition data sets are attractive to students occasionally, but I've seen a lot of them before and don't really want to see them again.

I cannot guarantee the quality of any of the data sets available at these sites, but I've spent at least a little time at many of them in recent months.

2.3.5 If you are storing your own data

An extremely useful link for those of you **building a spreadsheet to store data** is Karl Broman's tutorial on the subject. No one was born knowing this stuff - take a look.

2.4 Evaluating Task A

Dr. Love will evaluate all proposals (Task A) personally, in the order in which they are received. Proposals will receive one of two grades: OK or REDO. That grade will be posted to Canvas. REDO will be accompanied with specific requests in the form of a Canvas comment that should be accomplished within a short time window (approximately 24 hours). If you materially deviate from these specifications, Dr. Love will not evaluate your proposal other than to re-specify what needs to be fixed before he will respond.

- A score of OK is worth 10/10 points for Task A, once Task B is also complete.
- You (and/or your group, if there are problems with parts 1-3) must REDO the Proposal until you reach OK. Sometimes, that's more than once.

Chapter 3

Task B (Presentation Sign-Up) Instructions

3.1 Deadline and Submission information

Task B is due at noon on 2018-10-12. Submit your Task B work by completing the Google Form linked at **LINK GOES HERE**.

- Please note that Task A is also due at the same time.
- All students must specify a minimum of 8 time slots, on at least two different days, when they can give their presentation.
- You will also be able to specify your two favorite time slots among those you have chosen.
- The presentation dates are 2018-12-10, 2018-12-11 and 2018-12-13.
 - University classes end December 7.
 - December 10 is one of the official University Reading Days, and December 11 and 13 are Final Exam Days.
- If you have some special problem or concern or need to give your presentation before 2018-12-10, there is a space to tell Dr. Love about that at the end of the form.

Chapter 4

Task C (Survey Editing) Instructions

4.1 Deadline and Submission information

Task C is due at noon on 2018-10-23. Submit your Task C work via Canvas. Please note that:

- Task D is also due at the same time.
- We do not have class on 2018-10-23 because of CWRU's Fall Break.

Chapter 5

Task D (Survey Comparison Plan) Instructions

5.1 Deadline and Submission information

Task D is due at noon on 2018-10-23. Submit your Task D work via Canvas. Please note that:

- Task C is also due at the same time.
- We do not have class on 2018-10-23 because of CWRU's Fall Break.

Chapter 6

Task E (Taking the Survey) Instructions

6.1 Deadline and Submission information

Task E is due at noon on 2018-10-31. Submit your answers to the course survey via the Google Form linked at **LINK GOES HERE**. That link will go live after class on 2018-10-25.

Chapter 7

Task F (Sharing Study 2 Data) Instructions

7.1 Deadline and Submission information

Task F is due at noon on 2018-11-14. Submit your Task F work through Canvas.

7.2 Sharing Your Data Appropriately

Task F requires you to share your data for Study 2. The model for this Task is Jeff Leek's Guide to Data Sharing. Specifically, you will submit the following to Dr. Love by the deadline.

1. a direct link to the raw data set (without any need for me to sign up for anything) or a .csv copy of the raw data set called **yourname-raw.csv**
2. a single .csv file with a name of your choice containing a clean, tidy data set for Study 2, along with
3. a Word or PDF file containing both
 - a. a **codebook** section which describes every variable (column) and its values in your .csv file,
 - b. a **study design** section which reminds (and updates) us about the source of the data and your research question.

More to come.

Chapter 8

Task G (The Project Update) Instructions

8.1 Deadline and Submission information

Task G is due at noon on 2018-11-28. Submit your Task G work through Canvas.

Chapter 9

Task H (The Portfolio) Instructions

Task H requires you to provide a written portfolio of materials, which you will also make use of in your final presentation.

Details to come.

Chapter 10

Task I (Your Presentation) Instructions

Details to come.